

Ameya Dhamanaskar

Senior ML Engineer | Scalable AI Systems • Deep Learning • Production ML

adhamana@asu.edu | (949) 561-2038 | linkedin.com/in/ameyad95/ | github.com/nudlesoup

EXPERIENCE

Optum (UnitedHealth Group)

Machine Learning Engineer

Phoenix, Arizona

Jun 2023 - Present

- Architected AI optimization for pharmacy automation across 2,000+ NVIDIA Jetson devices, improving prescription throughput by 25% and reducing fill-time latency.
- Designed scalable computer vision pipelines combining synthetic and real-world datasets, maintaining >98% accuracy under variable lighting and motion.
- Evaluated and deployed real-time detection (YOLOv8) and tracking (DeepSORT) optimized with TensorRT and FP16/INT8 quantization, increasing inference speed by 40% and reducing memory by 30%.
- Containerized inference with Docker for reproducible deployment across Jetson environments, streamlining validation and rollout.
- Mentored 2 junior engineers and led design reviews that improved deployment reliability; drove cross-functional alignment with hardware, compliance, and operations on safety, regulatory, and performance KPIs.

Radius AI

Machine Learning Intern

Tempe, Arizona

Sep 2022 - May 2023

- Enhanced retail product detection using 3D sim2real transfer learning with RetinaNet and Faster R-CNN, achieving a +10 mAP gain while halving inference latency in edge deployments.
- Optimized Region Proposal Networks via pruning and mixed-precision inference, reducing model size and compute cost without accuracy loss.
- Leveraged Segment Anything Model (SAM) for automated labeling, doubling Vision Transformer (ViT) training efficiency and reducing manual annotation workload by 80%.

Institut de Robotica i Informatica Industrial (CSIC-UPC)

Machine Learning Researcher

Barcelona, Spain

Oct 2019 - Aug 2021

- Proposed a self-supervised egocentric 3D pose estimation framework combining first- and third-person cues; improved state-of-the-art accuracy by 12% on a 150k-frame dataset (Pattern Recognition, 2023).
- Built temporal CNN/Siamese architectures enforcing cross-view consistency for accurate 3D reconstruction from 2D keypoints.
- Modeled motion dynamics with encoder-decoder LSTMs to achieve robust 3D joint regression under head movement and occlusion.

Tesco Technologies

Software Development Engineer

Bangalore, India

Jun 2018 - Sep 2019

- Automated remote rebuilds for 50,000+ point-of-sale systems across the UK and Ireland; implemented Bi-Dijkstra and A*-based routing to reduce delivery times by 20%.

EDUCATION

Arizona State University

Masters (Hons.) in Computer Science, GPA: 4.0/4.0

Tempe, Arizona

Aug 2021 - Jun 2023

BITS Pilani

B.E (Hons.) in Electrical and Electronics Engineering

Rajasthan, India

Aug 2014 - May 2018

PUBLICATIONS

Enhancing Egocentric 3D Pose Estimation with Third Person Views - *Pattern Recognition*, 2023 | [\(DOI\)](#)

PROJECTS

Robustness Analysis of Object Detection Models

(PyTorch)

- Evaluated YOLOv5 and Mask R-CNN under noise/blur/transformations; findings guided production model selection and augmentation strategies.

Object Detection for Autonomous Vehicles

(PyTorch)

- Reimplemented Fast R-CNN, YOLOv3, and SSD with anchor-free detection and Soft-NMS; achieved +11% mAP on a KITTI subset.

TECHNICAL SKILLS

ML & CV

Object Detection, Tracking, Segmentation, 3D Pose Estimation, Transformers, Sim2Real.

Frameworks

PyTorch, TensorFlow, TensorRT, ONNX, OpenCV, CUDA, Quantization, Pruning.

MLOps & Deployment

Docker, CI/CD, Model Registry, Monitoring & Observability.

Programming

Python, C++ (Linux/Unix) | Familiar: Java, Matlab.